

Machine Learning Security Project



https://github.com/APruner-23/ML_Sec_Project.git



[Progetto_MLSec_Leandri_Pruner.ipynb](#)

Milena Leandri (70/90/00519) - Alessandro Pruner (70/90/00502)



01

Introduction

Objective of the Project

Objective of our project



Adversarial Generation

Generation of adversarial
examples on a 3 **RobustBench**
model **ensemble**



Transferability Evaluation

Evaluate Transferability on 7
different RobustBench models

Adversarial Crafting *characteristics*



Universal

Each adversarial example is crafted against the combined ensemble of three models



Untargeted

We do not specify any **target class**

We only want the models to **misclassify** the input sample



Model Ensemble



- Our **ensemble**:

| | Clean Accuracy | Robust Accuracy | Architecture |
|------------------------------------|----------------|-----------------|------------------|
| Zhang2019You | 87.20% | 44.83% | WideResNet-34-10 |
| Xu2023Exploring_WRN-28-10 | 93.69% | 63.89% | WideResNet-28-10 |
| Gowal2021Improving_28_10_ddpm_100m | 87.50% | 63.38% | WideResNet-28-10 |





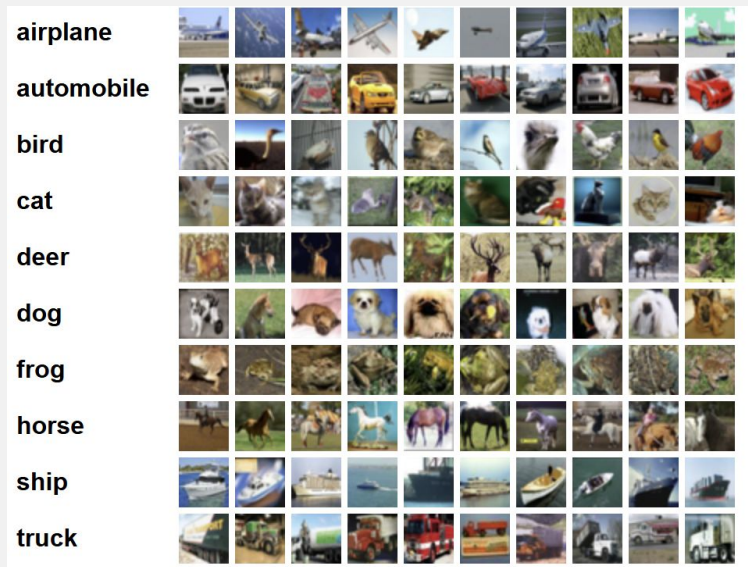
Transferability Models



| | Clean Accuracy | Robust Accuracy | Architecture |
|-------------------------------------|----------------|-----------------|----------------------------------|
| Huang2022Revisiting_WRN-A4 | 91.58% | 65.79% | WideResNet-A4 |
| Peng2023Robust | 93.27% | 71.07% | RaWideResNet-70-16 |
| Amini2024MeanSparse_Ra_WRN_70_16 | 93.24% | 68.94% | MeanSparse RaWideResNet-70-16 |
| Sehwag2021Proxy_ResNest152 | 87.30% | 62.79% | ResNet152 |
| Debenedetti2022Light_XCiT-L12 | 91.73% | 57.58% | XCiT-L12 |
| Cui2023Decoupled_WRN-28-10 | 92.16% | 67.73% | WideResNet-28-10 |
| Rebuffi2021Fixing_28_10_cutmix_ddpm | 87.33% | 60.73% | WideResNet-28-10 |

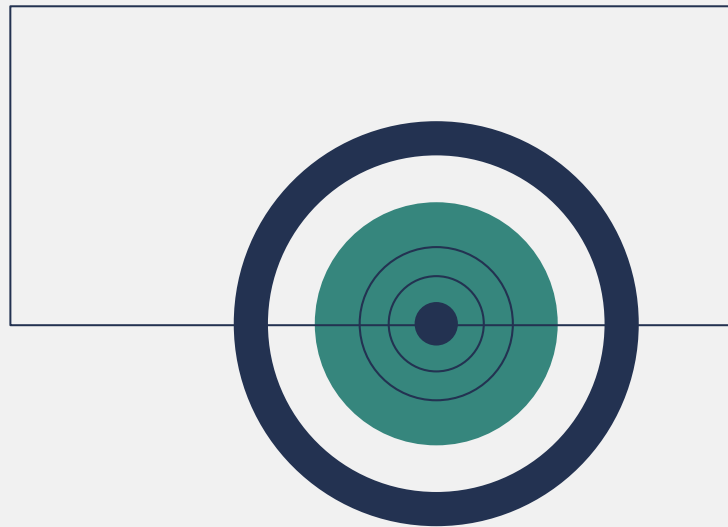
CIFAR-10 Dataset

- 60000 32x32 *colour* images
- 10 Classes
- 6000 images per class
- All *models* used have been **trained** on this dataset



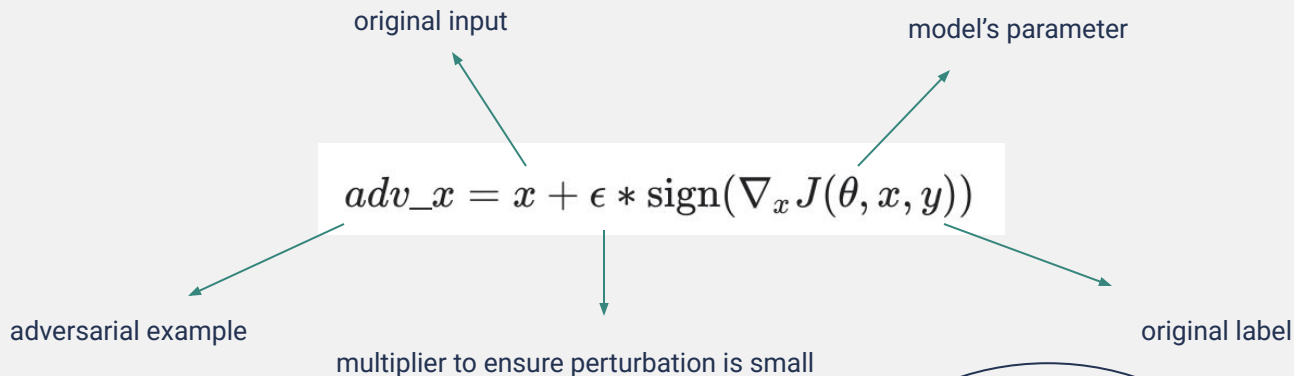
02

FGSM Attack



Fast Gradient Sign Method (FGSM)

- Introduced by Goodfellow et al. in *Explaining and Harnessing Adversarial Examples* (2015)
- Adversarial attack intended to reveal weaknesses in machine learning models
- Perturbs the input in the direction of the gradient of the loss with respect to the input



FGSM Ensemble

1. Prepare the adversarial input
2. Ensemble prediction
3. Loss and gradient calculation
4. FGSM perturbation
5. Clipping
6. Output

Attacking an ensemble of models makes the adversarial perturbation more effective and robust, as it considers the weaknesses of multiple networks at once.

```
def fgsm_ensemble(models, x, true_labels, epsilon):  
    x_adv = x.clone().detach()  
    x_adv.requires_grad = True  
    logits_list = [model(x_adv) for model in models]  
    stacked_logits = torch.stack(logits_list, dim=0)  
    ensemble_logits = torch.mean(stacked_logits, dim=0)  
    loss = margin_loss(ensemble_logits, true_labels)  
    grad = torch.autograd.grad(loss, x_adv)[0]  
    x_adv = x_adv.detach() + epsilon  
    * torch.sign(grad.detach())  
    x_adv = torch.clamp(x_adv, 0, 1).detach()  
    y_pred = predict_with_ensemble(models, x_adv)  
    y_pred_0 = y_pred[0]  
    return x_adv, y_pred_0
```



Attack

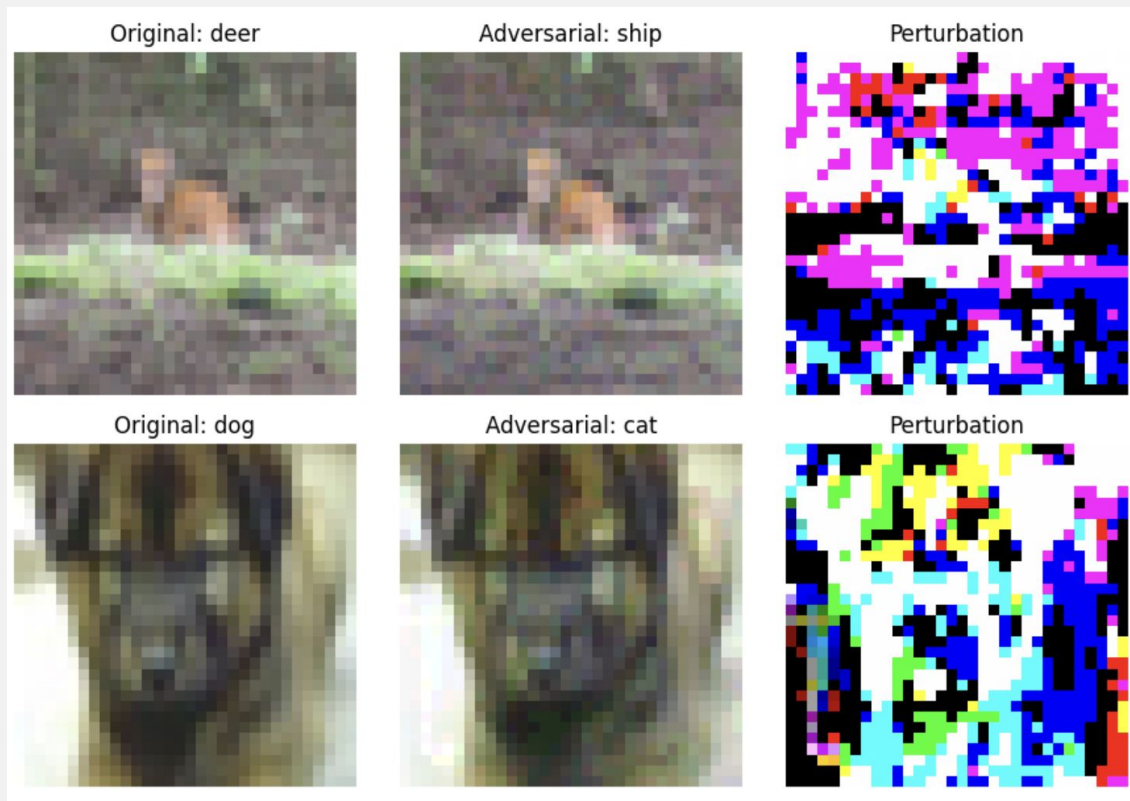
```
for (x, y_true) in data_loader:
    x, y_true = x.to(device), y_true.to(device)
    y_pred = predict_with_ensemble(models, x)
    y_pred_0 = y_pred[0]
    if y_pred_0 != y_true:
        excluded_adv_examples += 1
        continue
    x_adv, y_pred_adv = fgsm_ensemble(models, x, y_true,
8/255)
    if y_pred_adv != y_true:
        final_adversarials.append((x_adv, y_true.item()))
        total_adv_examples += 1
```

For each sample:

- **Predicts** the label using the *ensemble* of models
- If the ensemble already misclassifies the original input, that **sample is excluded** from the attack process
- Otherwise, *FGSM* adversarial attack is applied using an *epsilon* of 8/255 to **generate an adversarial example**



Generated Perturbation



Out of 1000 samples ...

149

adversarial examples created

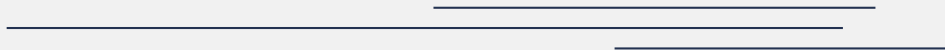
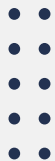
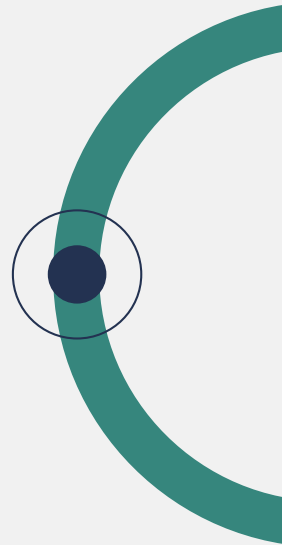
75

excluded examples



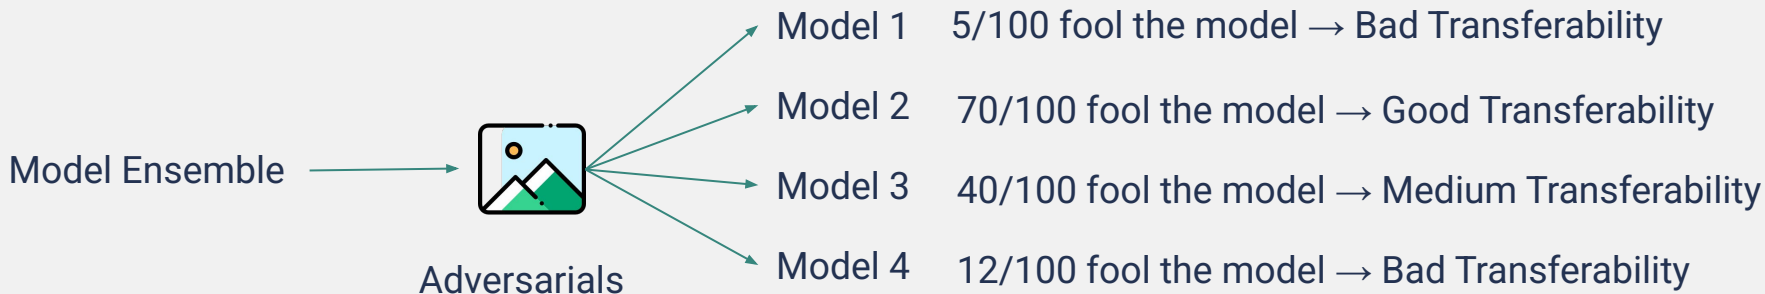
03

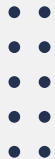
Transferability Evaluation



What's Transferability?

Transferability is useful to understand if an attack created to fool a model, also fools *other* models. It is the ability of an attack developed against a surrogate model to succeed also against a different target model

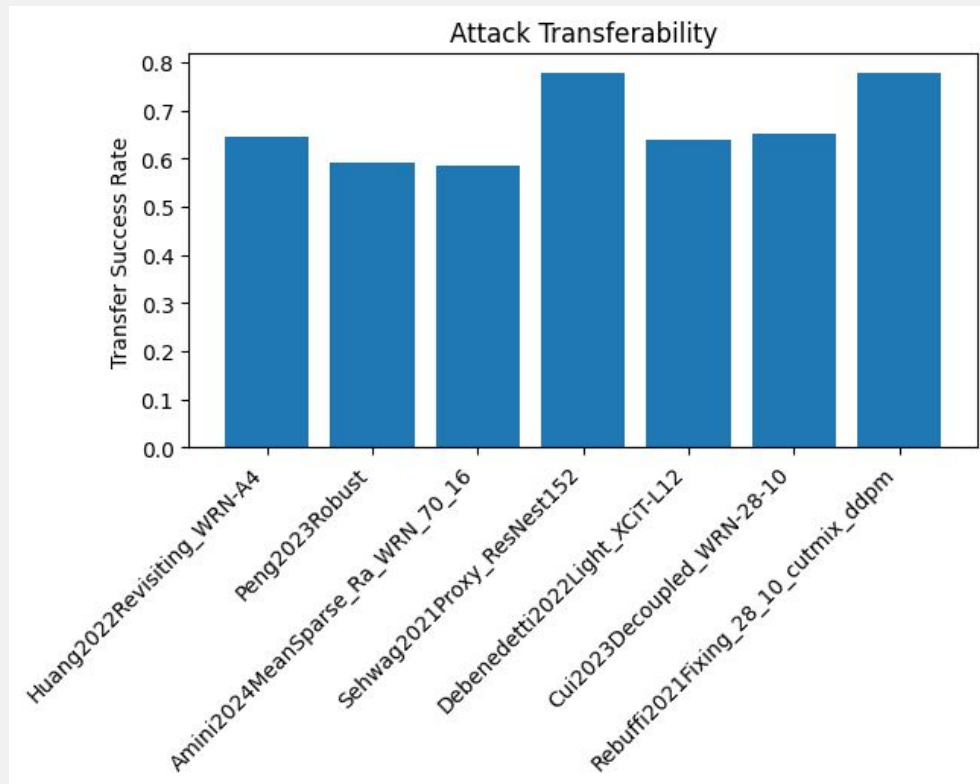




Results

| | Number of Adversarials that fooled the model | Attack Success Rate |
|---|--|---------------------|
| Huang2022Revisiting_WRN-A4 | 96 | 64.4% |
| Peng2023Robust | 88 | 59.1% |
| Amini2024MeanSparse_Ra_WRN _70_16 | 87 | 58.4% |
| Sehwag2021Proxy_ResNest152 | 116 | 77.9% |
| Debenedetti2022Light_XCiT-L12 | 95 | 63.8% |
| Cui2023Decoupled_WRN-28-10 | 97 | 65.1% |
| Rebuffi2021Fixing_28_10_cutmix _ddpm | 116 | 77.9% |

Attack Success Rate Plot





Thanks for your attention!

Milena Leandri (70/90/00519)
Alessandro Pruner (70/90/00502)

